

EE3233 Systems Programming for Engrs

Reference: M. Kerrisk, The Linux Programming Interface

Lecture 9 **Signals**



ECE ELECTRICAL & COMPUTER
ENGINEERING

What we will cover in this lecture


- Concept of Signal
- Types of Signal
- System Calls for Signal

Concepts

- Notification to a process that an event has occurred
 - Software interrupts
 - Interrupts the normal flow of execution of a program
 - One process can send a signal to another process: usually kernel sends it


Types of Events causing Kernel to Generate a Signal

- H/W exception: H/W detected a fault condition and then notifies kernel
 - dividing by 0
 - referencing a part of memory that is inaccessible
- User typed one of the terminal special characters that generate signals
 - interrupt character (*Control + C*)
 - suspend character (*Control + Z*)
- S/W event
 - terminal window resizing
 - child of this process terminated



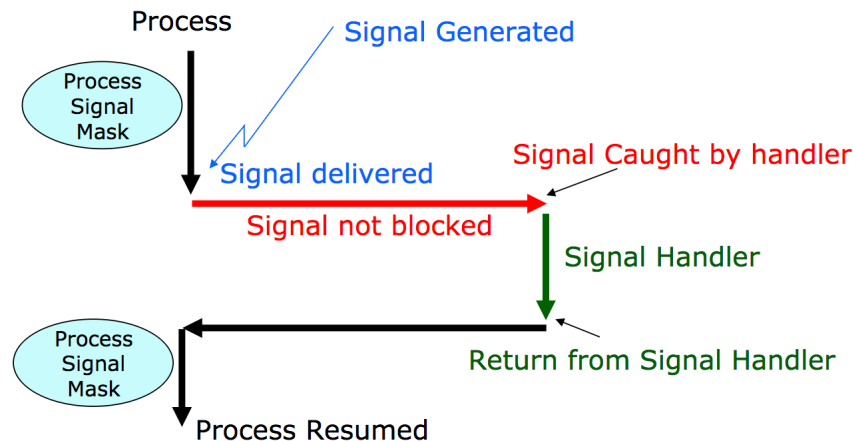
Concepts and Overview

- Each signal is defined as a unique integer
 - Starts from 1
 - Defined in `<signal.h>` with symbolic names of the form *SIGxxx*
 - When the user types the interrupt character, *SIGINT* (signal number **2**) is delivered to a process




Concepts and Overview

- Standard signal
 - Used by kernel to notify processes of event
 - On Linux, the standard signals are from 1 to 31
 - **Generated** by some event and later **delivered** to a process, which then takes some action in response to the signal
 - between '*generated*' and '*delivered*' times, a signal is said to be *pending*
- Normally, a pending signal is delivered to a process as soon as it is next scheduled to run, or immediately if the process is running



Concepts and Overview

- Sometimes, we need to ensure that a segment of code is not interrupted by the delivery of signal, then,
 - we can add a signal to the process's *signal mask*
- *Signal mask*: a set of signals whose delivery is currently blocked
 - If a signal is generated while it is blocked, it remains pending until it is later unblocked (removed from the signal mask)



Default Actions

- Signal is *ignored*
 - Discarded by the kernel and has no effect on the process
- The process is *terminated* (killed)
 - Abnormal process termination
 - Opposed to the normal process termination that occurs when a process terminates using *exit()*
- A *core dump file* is generated, and the process is terminated
 - Core dump file contains an image of the virtual memory of the process
- The process is *stopped*
 - Execution of the process is suspended
 - Execution of the process is resumed after previously being stopped

Disposition

- A program can change the action when the signal is delivered
 - Default action should occur: undo an earlier change
 - The signal is ignored
 - A signal handler is executed

Signal Handler

- A function, written by the programmer, that performs appropriate tasks in response to the delivery of a signal
- Some terminologies:
 - *installing* or *establishing*: notifying the kernel that a handler function should be invoked
 - *handled* or *caught*: when a signal handler is invoked in response to the delivery of a signal

Signal Types and Default Actions

- SIGABRT
 - A process is sent this signal when it calls the *abort()* function
 - By default, this signal terminates the process with a core dump intentionally for debugging
- SIGCHLD
 - is sent (by kernel) to a parent process when one of its children terminates, or stops or resumes by a signal
- SIGCONT
 - When sent to a stopped process, this signal causes the process to resume

Signal Types and Default Actions

- SIGFPE (Floating Point Error)
 - is generated for certain types of arithmetic errors, such as divide-by-zero
- SIGINT
 - When the user type the terminal interrupt character, the terminal driver sends this signal to the foreground process group
 - The default action for this signal is to terminate the process
- SIGKILL
 - *sure kill* signal
 - Can't be blocked, ignored, or caught by a handler, and thus always terminate a process

Signal Types and Default Actions

- SIGQUIT
 - When the user types the quit character (**Control + **) on the keyboard
 - This signal is sent to the foreground process group
 - By default, this signal terminates a process and causes it to produce a core dump for debugging
- SIGTERM
 - Terminates a process
 - Sent by the kill commands (Users sometimes explicitly kill a process using kill -9)

Changing Signal Disposition: *signal()*

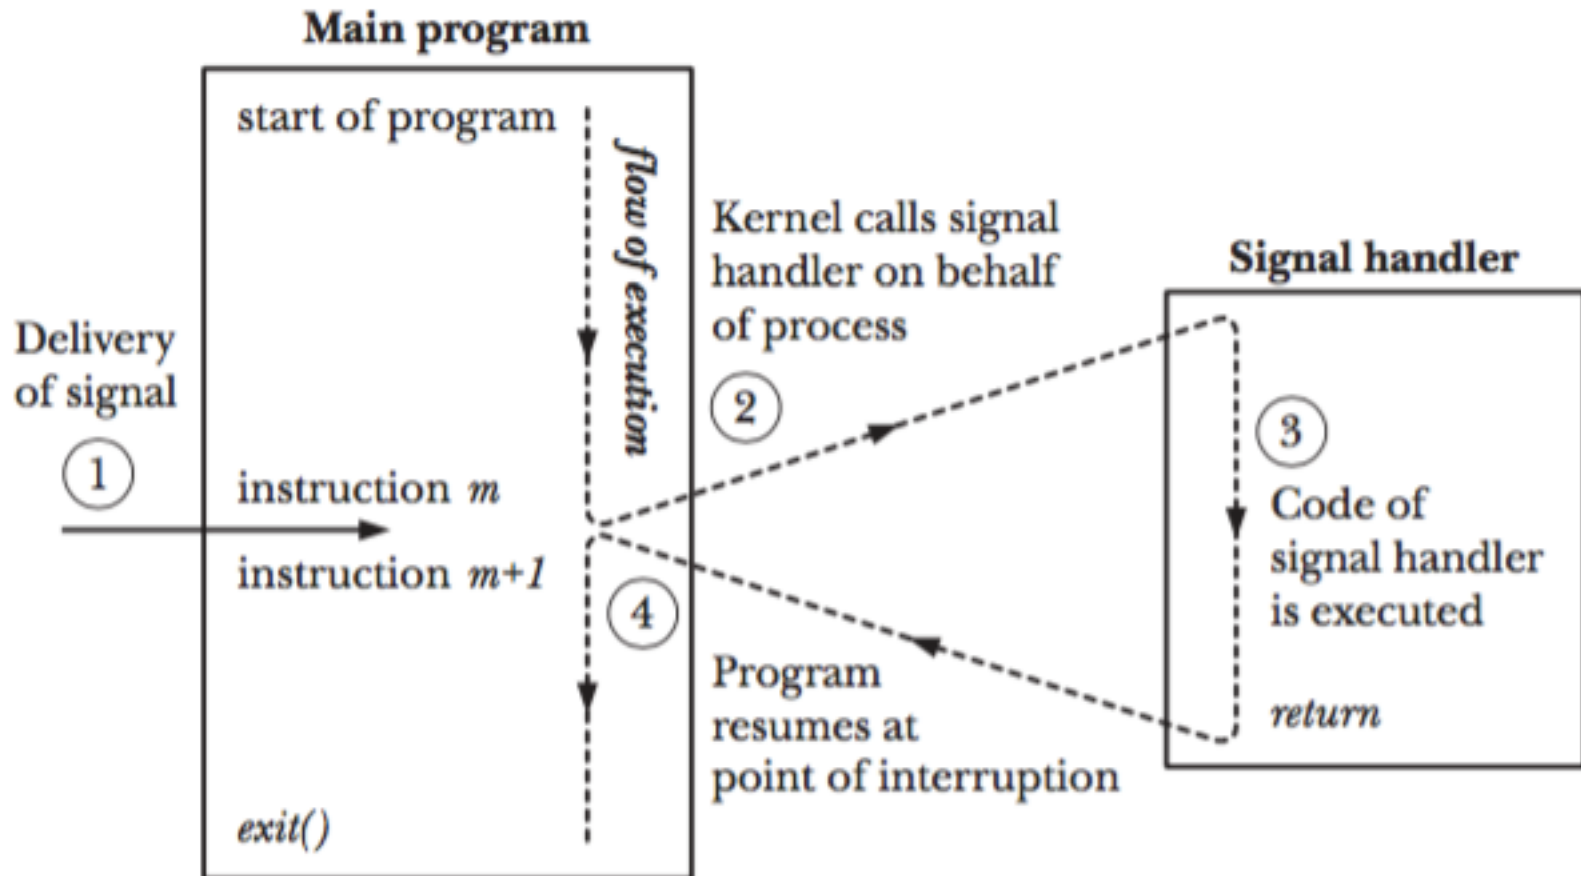
```
#include <signal.h>
```

```
void ( *signal(int sig, void (*handler)(int)) ) (int);
```

- The first argument, *sig* identifies the signal whose disposition we wish to change
- The second argument, *handler* is the address of the function that should be called when this signal is delivered

```
void handler (int sig) {  
    /* code for the handler */  
}
```

Changing Signal Disposition: *signal()*



Invocation of a signal handler may interrupt the main program flow at any time; the kernel calls the handler on the process's behalf, and when the handler returns, execution of the program resumes at the point where the handler interrupted it

Changing Signal Disposition: *signal()*

```
#include <signal.h>
#include "tlpi_hdr.h"

static void
sigHandler(int sig)
{
    printf("Ouch!\n");                /* UNSAFE (see Section 21.1.2) */
}

int
main(int argc, char *argv[])
{
    int j;

    if (signal(SIGINT, sigHandler) == SIG_ERR)
        errExit("signal");

    for (j = 0; ; j++) {
        printf("%d\n", j);
        sleep(3);                      /* Loop slowly... */
    }
}
```


Changing Signal Disposition: *signal()*

```
$ ./ouch
```

```
0
```

```
Type Control-C
```

```
Ouch!
```

```
1
```

```
2
```

```
Type Control-C
```

```
Ouch!
```

```
3
```

```
.
```

```
.
```

- The terminal driver generates *SIGINT* signal when we type terminal interrupt character (Control-C)
- The main program continuously loops
- When the kernel invokes a signal handler, it passes the number of the signal that caused the invocation as an integer to handler
- We can establish the same handler to catch different types of signals and use this argument to determine

Changing Signal Disposition: *signal()*

```
#include <signal.h>
#include "tlpi_hdr.h"

static void sigHandler(int sig) {
    static int count = 0;
    if (sig == SIGINT) {
        count++;
        printf("Caught SIGINT (%d)\n", count);
        return; /* Resume execution at point of interruption */
    }

    /* Must be SIGQUIT - print a message and terminate the process */
    printf("Caught SIGQUIT - that's all folks!\n");
    exit(EXIT_SUCCESS);
}

int main(int argc, char *argv[]) { /* Establish same handler for SIGINT and SIGQUIT */
    if (signal(SIGINT, sigHandler) == SIG_ERR) errExit("signal");
    if (signal(SIGQUIT, sigHandler) == SIG_ERR) errExit("signal");
    for (;;) /* Loop forever, waiting for signals */
        pause(); /* Block until a signal is caught */
}
```

Changing Signal Disposition: *signal()*

\$./intquit

^CCaught SIGINT(1)

^CCaught SIGINT(2)

^CCaught SIGINT(3)

^CCaught SIGINT(4)

^CCaught SIGINT(5)

^\Caught SIGQUIT – that’s all folks!

- The code of the handler distinguishes two signals by examining the *sig* argument, and takes different actions
- *pause()* blocks the process until a signal is caught

Sending Signals: kill()

- One process can send a signal to another process using the **kill()** system call

```
#include <signal.h>
```

```
int kill(pid_t pid, int sig);
```


- *pid* identifies one or more processes to which the signal specified by *sig* is to be sent
- Returns 0 on success, or -1 on error

Sending Signals: *kill()*

- ***pid* > 0**
 - the signal is sent to the process with the process ID specified by *pid*
- ***pid* == 0**
 - the signal is sent to every process in the same process group as the calling process (including itself)
- ***pid* < -1**
 - the signal is sent to all of the processes in the process group whose ID equals the absolute value of *pid*
- ***pid* == -1**
 - the signal is sent to every process for which the calling process has permission to send a signal, except *init* and the calling process

Sending Signals: *kill()*

- A privileged (CAP_KILL) process may send a signal to any process
- The *init* process (PID 1) can be sent only signals for which it has a handler installed
 - this prevents from accidentally killing *init*
- An unprivileged process can send a signal to another process if
 - real or effective UID of sending process matches real or saved UID of receiving process



Checking existence of a process

- *kill()* can serve to check existence of a process
- If the *sig* argument is specified as 0, then no signal is sent
 - Instead, merely performs error checking to see if the process can be signaled
 - If sending a null signal fails with the error *ESRCH*, then we know the process doesn't exist
 - If fails with the error *EPERM* (meaning we do not have permission to send a signal to the process) or succeeds, the process exists

```

#include <signal.h>
#include "tlpi_hdr.h"

int
main(int argc, char *argv[])
{
    int s, sig;

    if (argc != 3 || strcmp(argv[1], "--help") == 0)
        usageErr("%s sig-num pid\n", argv[0]);

    sig = getInt(argv[2], 0, "sig-num");

    s = kill(getLong(argv[1], 0, "pid"), sig);

    if (sig != 0) {
        if (s == -1)
            errExit("kill"); /* Something wrong */

    } else { /* Null signal: process existence check */
        if (s == 0) {
            printf("Process exists and we can send it a signal\n");
        } else {
            if (errno == EPERM)
                printf("Process exists, but we don't have "
                    "permission to send it a signal\n");
            else if (errno == ESRCH)
                printf("Process does not exist\n");
            else
                errExit("kill"); /* Something wrong */
        }
    }

    exit(EXIT_SUCCESS);
}

```


Checking existence of a process

```
$ ./ouch
```

```
1
```

```
2
```

```
3
```

```
4
```

```
5
```

```
*
```

```
*
```

```
*
```

```
$ ps aux | grep ouch  
user1 3429 ... ./ouch
```

```
$ ./t_kill 3429 0
```

Process exists and we
can send it a signal

```
$ ./t_kill 3333 0
```

Process does not exist

Other Ways of Sending Signals: *raise()*

- Sometimes useful to send a signal to itself

```
#include <signal.h>
```

```
int raise(int sig);
```

- In a single-threaded program, a call to *raise()* is equivalent to the following call to *kill()*

kill(getpid(), sig);

- When a process sends itself a signal using *raise()*, the signal is delivered immediately

Other Ways of Sending Signals: *killpg()*

- Sometimes useful to send a signal to itself

```
#include <signal.h>
```

```
int killpg(pid_t pgrp, int sig);
```

- equivalent to the following call to `kill()`:

`kill(-pgrp, sig);`

If *pgrp* is specified as 0, then the signal is sent to all processes in the same process group as the caller

Displaying Signal Descriptions

- `char *strsignal (int sig);`
 - Each signal has an associated printable description: These descriptions are listed in the array *sys_siglist* – For example, we can refer to *sys_siglist[SIGPIPE]* to get the description for *SIGPIPE* (broken pipe)
 - performs bounds checking on the sig argument, and then returns a pointer to a printable description of the signal

Signal Sets

- **int sigemptyset (sigset_t *set);**
 - initializes a signal set to contain no members
- **int sigfillset (sigset_t *set);**
 - initializes a signal set to contain all signals
- **int sigaddset (const sigset_t *set, int sig);**
 - individual signals can be added to a *set*
- **int sigdelset (const sigset_t *set, int sig);**
 - individual signals can be removed
 - returns true if *set* contains no signals

Signal Sets

- **int sigisemptyset (const sigset_t *set);**
 - returns true if *set* contains no signals
- **int sigismember (const sigset_t *set, int sig);**
 - returns 1 if *sig* is a member of *set*, and 0 otherwise
- **int sigandset (sigset_t *dest, sigset_t *left, sigset_t *right);**
 - places the intersection of the sets *left* and *right* in the set *dest*
- **int sigorset (sigset_t *dest, sigset_t *left, sigset_t *right);**
 - places the union of the sets *left* and *right* in the set *dest*

Signal Mask (Blocking Signal Delivery)

- **int sigprocmask (int *how*, const sigset_t **set*, sigset_t **oldset*);**
 - used to explicitly add signals to and remove signals from the signal mask
 - *how*:
 - SIG_BLOCK**: *set* is added to the signal mask (union)
 - SIG_UNBLOCK**: *set* is removed from the signal mask
 - SIG_SETMASK**: *set* is assigned to the signal mask
 - *oldset*: if the *oldset* is not NULL, it points to a *sigset_t* buffer that is used to return the previous signal mask

Signal Mask (Blocking Signal Delivery)

```
sigset_t blockSet, prevMask;

/* Initialize a signal set to contain SIGINT */

sigemptyset(&blockSet);
sigaddset(&blockSet, SIGINT);

/* Block SIGINT, save previous signal mask */
if (sigprocmask(SIG_BLOCK, &blockSet, &prevMask) == -1)
    errExit("sigprocmask1");

/* ... Code that should not be interrupted by SIGINT ... */

/* Restore previous signal mask, unblocking SIGINT */
if (sigprocmask(SIG_SETMASK, &prevMask, NULL) == -1)
    errExit("sigprocmask2");
```

(1)

(2)

Signal mask

Signal mask

- Temporarily prevents delivery of a signal(SIGINT) and then unblock it by resetting the signal mask to its previous state

Pending Signals

- **int sigpending (sigset_t *set);**
 - If a process receives a signal that it is currently blocking, that signal is added to the process's set of pending signals
 - *sigpending()* system call returns the set of signals that are pending for the calling process in the *sigset_t* structure pointed to by *set*

Example Program Displaying Signal Sets

```
void                                /* Print list of signals within a signal set */
printSigset(FILE *of, const char *prefix, const sigset_t *sigset)
{
    int sig, cnt;

    cnt = 0;
    for (sig = 1; sig < NSIG; sig++) {
        if (sigismember(sigset, sig)) {
            cnt++;
            fprintf(of, "%s%d (%s)\n", prefix, sig, strsignal(sig));
        }
    }

    if (cnt == 0)
        fprintf(of, "%s<empty signal set>\n", prefix);
}
```

Example Program Displaying Signal Sets

```
int                                     /* Print mask of blocked signals for this process */
printStatsMask(FILE *of, const char *msg)
{
    sigset_t currMask;

    if (msg != NULL)
        fprintf(of, "%s", msg);

    if (sigprocmask(SIG_BLOCK, NULL, &currMask) == -1)
        return -1;

    printSigset(of, "\t\t", &currMask);

    return 0;
}
```

Example Program Displaying Signal Sets

```
int          /* Print signals currently pending for this process */
printPendingSigs(FILE *of, const char *msg)
{
    sigset_t pendingSigs;

    if (msg != NULL)
        fprintf(of, "%s", msg);

    if (sigpending(&pendingSigs) == -1)
        return -1;

    printSigset(of, "\t\t", &pendingSigs);

    return 0;
}
```

Signals are not queued

signals/sig_sender.c

```
#include <signal.h>
#include "tlpi_hdr.h"

int
main(int argc, char *argv[])
{
    int numSigs, sig, j;
    pid_t pid;

    if (argc < 4 || strcmp(argv[1], "--help") == 0)
        usageErr("%s pid num-sigs sig-num [sig-num-2]\n", argv[0]);

    pid = getLong(argv[1], 0, "PID");
    numSigs = getInt(argv[2], GN_GT_0, "num-sigs");
    sig = getInt(argv[3], 0, "sig-num");
```

Signals are not queued

```
/* Send signals to receiver */

printf("%s: sending signal %d to process %ld %d times\n",
       argv[0], sig, (long) pid, numSigs);

for (j = 0; j < numSigs; j++)
    if (kill(pid, sig) == -1)
        errExit("kill");

/* If a fourth command-line argument was specified, send that signal */

if (argc > 4)
    if (kill(pid, getInt(argv[4], 0, "sig-num-2")) == -1)
        errExit("kill");

printf("%s: exiting\n", argv[0]);
exit(EXIT_SUCCESS);
}
```

Signal receiver

```
#define _GNU_SOURCE
#include <signal.h>
#include "signal_functions.h"          /* Declaration of printSigset() */
#include "tlpi_hdr.h"

static int sigCnt[NSIG];              /* Counts deliveries of each signal */
static volatile sig_atomic_t gotSigint = 0;
                                      /* Set nonzero if SIGINT is delivered */

static void
handler(int sig)
{
    if (sig == SIGINT)
        gotSigint = 1;
    else
        sigCnt[sig]++;
}

int
main(int argc, char *argv[])
{
    int n, numSecs;
    sigset_t pendingMask, blockingMask, emptyMask;

    printf("%s: PID is %ld\n", argv[0], (long) getpid());
```

```

for (n = 1; n < NSIG; n++)          /* Same handler for all signals */
    (void) signal(n, handler);      /* Ignore errors */

/* If a sleep time was specified, temporarily block all signals,
   sleep (while another process sends us signals), and then
   display the mask of pending signals and unblock all signals */

if (argc > 1) {
    numSecs = getInt(argv[1], GN_GT_0, NULL);

    sigfillset(&blockingMask);
    if (sigprocmask(SIG_SETMASK, &blockingMask, NULL) == -1)
        errExit("sigprocmask");

    printf("%s: sleeping for %d seconds\n", argv[0], numSecs);
    sleep(numSecs);

    if (sigpending(&pendingMask) == -1)
        errExit("sigpending");

    printf("%s: pending signals are: \n", argv[0]);
    printSigset(stdout, "\t\t", &pendingMask);

    sigemptyset(&emptyMask);        /* Unblock all signals */
    if (sigprocmask(SIG_SETMASK, &emptyMask, NULL) == -1)
        errExit("sigprocmask");
}

```


Signal receiver

```
while (!gotSigint)                /* Loop until SIGINT caught */
    continue;

for (n = 1; n < NSIG; n++)        /* Display number of signals received */
    if (sigCnt[n] != 0)
        printf("%s: signal %d caught %d time%s\n", argv[0], n,
                sigCnt[n], (sigCnt[n] == 1) ? "" : "s");

exit(EXIT_SUCCESS);
}
```

Signal receiver

```
$ ./sig_receiver 15 &                                Receiver blocks signals for 15 secs
[1] 5368
./sig_receiver: PID is 5368
./sig_receiver: sleeping for 15 seconds

$ ./sig_sender 5368 1000000 10 2                    Send SIGUSR1 signals, plus a SIGINT
./sig_sender: sending signal 10 to process 5368 1000000 times
./sig_sender: exiting
./sig_receiver: pending signals are:
                2 (Interrupt)
                10 (User defined signal 1)
./sig_receiver: signal 10 caught 1 time
[1]+  Done                                     ./sig_receiver 15
```

A blocked signal is delivered only once, no matter how many times it is generated